

Technical Specification Document

Project: Quick Best-Performance RAG System for Systematic Review

Name: Clifford Hepplethwaite

Email: clifford@tumpetech.com

1. Overview & Purpose

This project implements a high-performance Retrieval-Augmented Generation (RAG) system designed to support systematic review workflows by enabling researchers to query academic PDFs and receive accurate, context-aware answers with source traceability.

2. Architecture Overview

```
graph TD
  A[PDF Documents] --> B[LlamaIndex PDFLoader]
  B --> C[SectionNodeParser Chunking]
  C --> D[OpenAI Embeddings (text-embedding-3-large)]
  D --> E[Qdrant Cloud Vector Store]
  F[User Query] --> G[LlamaIndex Retriever (+ optional reranker)]
  E --> G
  G --> H[OpenAI GPT-4 Turbo LLM]
  H --> I[Streamlit Frontend]
```

3. Data Sources & Formats

- Input data:** Academic PDF documents (systematic review papers, reports)
- Document ingestion:** Use LlamaIndex's built-in PDFLoader which extracts text and metadata (title, author, etc.)
- Chunking strategy:** Use SectionNodeParser with overlapping chunks to preserve document structure and context boundaries (e.g., Introduction, Methods, Results)

4. Technology Stack

Layer	Tool & Setup	Notes
Ingestion	LlamaIndex with PDFLoader	Install via <code>pip install llama-index</code>
Chunking	SectionNodeParser with overlap	Preserve academic structure for better QA

Layer	Tool & Setup	Notes
Embeddings	OpenAI text-embedding-3-large	Requires OpenAI API key; simple Python API call
Vector Store	Qdrant Cloud	Managed service; free tier available; API based
Retriever	LlamaIndex hybrid retriever (+ optional reranker)	Hybrid semantic + lexical search; reranker optional
LLM	OpenAI GPT-4 Turbo	Fast, high-quality text generation; requires API key
Frontend	Streamlit	<code>pip install streamlit</code> ; minimal UI for querying

5. System Components Details

5.1 Document Ingestion & Processing

- PDFs are loaded via LlamaIndex PDFLoader.
- Extract text and metadata for each document.
- Store documents in memory for chunking.

5.2 Chunking

- Use `SectionNodeParser` to split documents by sections.
- Overlap chunks by ~20% to maintain context across boundaries.

5.3 Embedding Generation

- Use OpenAI's embedding API `text-embedding-3-large` for semantic vector creation.
- Each chunk is embedded and stored in Qdrant.

5.4 Vector Storage & Retrieval

- Qdrant Cloud stores all vectors and supports similarity search.
- Hybrid retrieval combines vector similarity with lexical search via LlamaIndex retriever.

5.5 LLM Integration

- OpenAI GPT-4 Turbo used for answer generation.
- Queries combine retrieved context chunks into prompt with instructions for concise, citation-aware answers.

5.6 Frontend

- Streamlit app to accept user queries and display answers.
 - Supports file upload for new PDFs and real-time query processing.
-

6. Integration Points

- OpenAI API for embeddings and GPT-4 Turbo calls.
 - Qdrant Cloud API for vector indexing and similarity search.
 - Streamlit web server for frontend.
-

7. Performance & Scalability Considerations

- Use batching for embeddings to reduce API calls.
 - Qdrant indexes scale to millions of vectors with low latency.
 - GPT-4 Turbo supports faster generation than base GPT-4, suitable for real-time query handling.
 - Frontend designed for single-user or small team use; scale backend with FastAPI + React if needed.
-

8. Error Handling & Logging

- Handle API rate limits and retries for OpenAI calls.
 - Validate PDF uploads for compatibility.
 - Log user queries and errors with timestamp.
 - Provide user-friendly error messages on UI.
-

9. Security & Privacy

- Store API keys securely using environment variables (`.env`).
 - Transmit data over HTTPS to Qdrant and OpenAI.
 - No user documents persist beyond session unless explicitly saved.
-

10. Testing Plan

- Unit test ingestion and chunking on sample PDFs.
 - Integration test embedding generation and vector storage.
 - End-to-end test query-answer flow in Streamlit UI.
 - Test with multiple document types and large PDFs.
-

11. Deployment Plan

- Local development with Python 3.9+ environment.
 - Deploy frontend on Streamlit Cloud or similar PaaS.
 - Qdrant Cloud handles vector DB hosting.
 - OpenAI API keys configured via environment variables.
-

12. Maintenance & Support

- Update dependencies quarterly.
 - Monitor API usage and costs monthly.
 - Expand to reranker model (e.g. Cohere or bge-reranker) after MVP launch.
 - Add user authentication if needed for multi-user scenarios.
-

Appendix

Sample Environment Variables

```
OPENAI_API_KEY=your_openai_api_key_here  
QDRANT_API_KEY=your_qdrant_api_key_here
```

Sample CLI Commands

```
pip install llama-index streamlit openai qdrant-client  
streamlit run app.py
```
